

A Major Project Report on
DeepTrace: A Deep Learning Framework for
Detection of AI-Generated and Deepfake Images
Using Hybrid Spatial–Frequency Domain Feature
Analysis

Submitted in Partial Fulfillment to Osmania University for the Award of the Degree
of

Bachelor of Engineering

In

Department of CSE (Data Science)

By

Mohammed Safiuddin

160922750037

Under the Esteemed Supervision of

Mr. Aqeel Haider Shams

Assistant Professor

Dept. of CSE – Data Science



DEPARTMENT OF CSE (DATA SCIENCE)

LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

UGC Autonomous

Approved by AICTE | Affiliated to Osmania University | Estd.2003
Sy.No.32, Himayat Sagar, Near TGPA Junction, Hyderabad-500091, India.

2025- 2026



DEPARTMENT OF CSE (DATA SCIENCE)
LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY
UGC Autonomous

Approved by AICTE | Affiliated to Osmania University | Estd.2003
Sy.No.32, Himayat Sagar, Near TGPA Junction, Hyderabad-500091, India.

Certificate

This is to certify that Major Project Report entitled “**DeepTrace: A Deep Learning Framework for Detection of AI-Generated and Deepfake Images Using Hybrid Spatial–Frequency Domain Feature Analysis**” is a bonafide record of the work successfully completed and submitted by

Mohammed Safiuddin

160922750037

Under the supervision of **Mr. Aqeel Haider Shams, Assistant Professor Department** of CSE (Data Science) for the requirement of partial fulfillment for the award of degree of Bachelor of Engineering in CSE (Data Science) during the academic year 2022-2026 from Osmania University.

The result embodied in this project report has not been submitted to any other university/institute for the award of any degree.

Project Guide

Head of the Department

Principal

External

Declaration

We hereby declare that Major Project entitled “**DeepTrace: A Deep Learning Framework for Detection of AI-Generated and Deepfake Images Using Hybrid Spatial–Frequency Domain Feature Analysis**” is being submitted by us in partial fulfilment to Osmania university for the award of Bachelor of Engineering in the Department of CSE (Data Science) at the Lords Institute of Engineering and Technology Autonomous, Hyderabad, Telangana, is the result of investigations carried out by us under the Guidance of **Mr. Syed Mushtaq Ali, Assistant Professor**, Department of CSE (Data Science), Lords Institute of Engineering and Technology.

The work is original and has not been submitted for any Degree or Diploma for this or any other university.

Mohammed Safiuddin

160922750037

Acknowledgement



MOHAMMED SAFIUDDIN

160922750037

9642580952

mohdsafiuddin99@gmail.com

In the name of almighty, the most beneficent and the most merciful, we thank the lord for helping us in all the stages of this thesis work.

We would like to thank our guide **Mr.Syed Mushtaq Ali, Assistant Professor and HoD** for his aspiring guidance and friendly advice during the project work. He was a constant motivation throughout the internship and took out some of his valuable time whenever it was required for which we are very grateful to him. We would be grateful for his motivation and co-operation at various stages, which provided us the knowledge base and the understanding of the potential internship in this field finally. We also appreciate the concern of my parents, friends and other staff members at the college.

We would like to thanks **Mr. Syed Mushtaq Ali, Head of the Department** for his support throughout our Project. It gives us great pleasure in acknowledging the support and help from him throughout the project. Our experience here has been full of opportunities to learn from faculty with a wide and deep expertise in Engineering.

We would like to thank our beloved **Principal Dr. Ravi Kishore Singh**. And we would like to directly thank those people in the department who was particularly instrumental in contributing to our experience at Lords Institute of Engineering and Technology.

We thank both teaching and non-teaching staff members of CSE (Data Science) department for their kind cooperation and all sorts of help to bring out this project work successfully.

Lastly, we derive great pleasure and we would like to specially thank our parents for their patience, continuous encouragement, support and love given towards completion of our research work.



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

(UGC Autonomous)

Approved by AICTE | Affiliated to Osmania University | Estd.2003.

Accredited with 'A' grade by NAAC | Accredited by NBA

Vision of the Institute.

Lords Institute of Engineering and Technology strives for excellence in professional education through quality, innovation and teamwork and aims to emerge as a premier institute in the state and across the nation.

Mission of the Institute.


- To impart quality professional education that meets the needs of present and emerging technological world.
- To strive for student achievement and success, preparing them for life, career and leadership.
- To provide a scholarly and vibrant learning environment that enables faculty, staff and students to achieve personal and professional growth.
- To contribute to advancement of knowledge, in both fundamental and applied areas of engineering and technology.
- To forge mutually beneficial relationships with government organizations, industries, society and the alumni.

PRINCIPAL

Lords Institute of Engineering & Tech.

(An Autonomous Institution)

Sy. No. 32, Himayath Sagar, Hyderabad-500091.T.S.



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

(UGC Autonomous)

Approved by AICTE | Affiliated to Osmania University | Estd.2003.

Accredited with 'A' grade by NAAC | Accredited by NBA

Department of CSE (Data Science)

Vision of the Department.

To emerge as center of excellence in Data Science education and innovation addressing the societal and industrial needs.

Mission of the Department.

DM1: Provide rigorous training and skills in Data Science. Including Artificial Intelligence, Machine Learning, Big Data Analytics, and Data Visualization.

DM2: Promote interdisciplinary research and collaboration in the application of Data Science.


DM3: Collaborate with industry and society towards overall development of stakeholders.

DM4: Practice ethical, contemporary tools, and self-learning to contribute responsibly and effectively in a diverse professional environment.

Note. DM. Department Mission


Head of the Department

Head of the Department
CSE (Data Science)
Lords Institute of Engg. & Tech.
Hyderabad-500091 T.S.



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

(UGC Autonomous)

Approved by AICTE | Affiliated to Osmania University | Estd.2003.

Accredited with 'A' grade by NAAC | Accredited by NBA


Department of CSE (Data Science)

B.E. CSE (Data science) Program Educational Objectives PEOs.

PEO1	Exhibit strong foundation in Basic Science, Mathematics, Statistics, Data Science, Computer Science and Allied engineering.
PEO2	Develop Data Science applications at advanced level for employability in IT industry especially in niche technologies.
PEO3	Demonstrate ethical responsibilities & commitment to teamwork and leadership.
PEO4	Apply interdisciplinary Data Science knowledge to develop innovative & sustainable solutions.


Head of the Department

Head of the Department
CSE (Data Science)
Lords Institute of Engg. & Tech.
Hyderabad-500091 T.S.



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

(UGC Autonomous)

Approved by AICTE | Affiliated to Osmania University | Estd.2003.

Accredited with 'A' grade by NAAC | Accredited by NBA

Department of CSE (Data Science)

B.E. CSE (Data science) Program Outcomes POs. Engineering

Graduates will be able to.

S. No.	Program Outcomes POs
1.	PO1: Engineering Knowledge: Apply knowledge of mathematics, natural science, computing, engineering fundamentals, and an engineering specialization as specified in WK1 to WK4 respectively to develop to the solution of complex engineering problems.
2.	PO2: Problem Analysis: Identify, formulate, review research literature, and analyse complex engineering problems reaching substantiated conclusions with consideration for sustainable development. (WK1 to WK4)
3.	PO3: Design/Development of Solutions: Design creative solutions for complex engineering problems and design/develop systems/components/processes to meet identified needs with consideration for the public health and safety, whole-life cost, net zero carbon, culture, society, and environment as required. (WK5)
4.	PO4: Conduct Investigations of Complex Problems: Conduct investigations of complex engineering problems using research-based knowledge including design of experiments, modelling, analysis & interpretation of data to provide valid conclusions. (WK8).
5.	PO5: Engineering Tool Usage: Create, select, and apply appropriate techniques, resources and modern engineering & IT tools, including prediction and modelling recognizing their limitations to solve complex engineering problems. (WK2 and WK6)
6.	PO6: The Engineer and The World: Analyse and evaluate societal and environmental aspects while solving complex engineering problems for its impact on sustainability with reference to economy, health, safety, legal framework, culture, and environment. (WK1, WK5, and WK7).
7.	PO7: Ethics: Apply ethical principles and commit to professional ethics, human values, diversity, and inclusion adhere to national & international laws. (WK9)
8.	PO8: Individual and Collaborative Team work: Function effectively as an individual, and as a member or leader in diverse/multi-disciplinary teams.
9.	PO9: Communication: Communicate effectively and inclusively within the engineering community and society at large, such as being able to comprehend and write effective reports and design documentation, make effective presentations considering cultural, language, and learning differences

10.	PO10: Project Management and Finance: Apply knowledge and understanding of engineering management principles and economic decision-making and apply these to one's own work, as a member and leader in a team, and to manage projects and in multidisciplinary environments.
11.	PO11: Life-Long Learning: Recognize the need for, and have the preparation and ability for i) independent and life-long learning ii) adaptability to new and emerging technologies and iii) critical thinking in the broadest context of technological change. (WK8)

Knowledge and Attitude Profile (WK)


WK's	Description
WK1	A systematic, theory-based understanding of the natural sciences applicable to the discipline and awareness of relevant social sciences.
WK2	Conceptually-based mathematics, numerical analysis, data analysis, statistics and formal aspects of computer and information science to support detailed analysis and modelling applicable to the discipline.
WK3	A systematic, theory-based formulation of engineering fundamentals required in the engineering discipline.
WK4	Engineering specialist knowledge that provides theoretical frameworks and bodies of knowledge for the accepted practice areas in the engineering discipline; much is at the forefront of the discipline.
WK5	Knowledge, including efficient resource use, environmental impacts, whole-life cost, re-use of resources, net zero carbon, and similar concepts, that supports engineering design and operations in a practice area.
WK6	Knowledge of engineering practice (technology) in the practice areas in the engineering discipline.
WK7	Knowledge of the role of engineering in society and identified issues in engineering practice in the discipline, such as the professional responsibility of an engineer to public safety and sustainable development.
WK8	Engagement with selected knowledge in the current research literature of the discipline, awareness of the power of critical thinking and creative approaches to evaluate emerging issues.
WK9	Ethics, inclusive behavior and conduct. Knowledge of professional ethics, responsibilities, and norms of engineering practice. Awareness of the need for diversity by reason of ethnicity, gender, age, physical ability etc. with mutual understanding and respect, and of inclusive attitudes.

Program Outcome	WK1	WK2	WK3	WK4	WK5	WK6	WK7	WK8	WK9
PO1 Engineering Knowledge	■	■	■	■					
PO2 Problem Analysis	■	■	■	■					

PO3 Design/Development of Solutions					■				
PO4 Conduct Investigations of Complex Problems								■	
PO5 Engineering Tool Usage		■				■			
PO6 The Engineer & The World	■				■		■		
PO7 Ethics									■
PO8 Individual & Collaborative Teamwork									
PO9 Communication									
PO10 Project Management and Finance									
PO11 Life-Long Learning								■	


Head of the Department

Head of the Department
CSE (Data Science)
Lords Institute of Engg. & Tech.
Hyderabad-500091 T.S.



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

(UGC Autonomous)

Approved by AICTE | Affiliated to Osmania University | Estd.2003.

Accredited with 'A' grade by NAAC | Accredited by NBA

Department of CSE (Data Science)

B.E. CSE (Data science) Program Specific Outcomes PSO's.

PSO1	Professional Skills: Apply Data Science in real-time software project development using programming environment to deliver quality software product.
PSO2	Problem-Solving Skills: Demonstrate ethical responsibility, data privacy awareness & effective communication delivering sustainable & socially relevant solutions


Head of the Department

Head of the Department
CSE (Data Science)
Lords Institute of Engg. & Tech.
Hyderabad-500091 T.S.



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

(UGC Autonomous)

Approved by AICTE | Affiliated to Osmania University | Estd.2003.


Accredited with 'A' grade by NAAC | Accredited by NBA

Department of CSE (Data Science)

Course Outcomes. C85 - Major Project

Student will be able to

CO. No	Description	Blooms Taxonomy Level
C85.1	Understand the core concepts of academic management systems, student-teacher interactions, and collaboration tools.	BTL2
C85.2	Apply software development techniques to design and implement features like attendance, grading, and communication modules.	BTL3
C85.3	Analyse system requirements and optimize workflows for better academic collaboration and efficiency.	BTL4
C85.4	Design and develop a complete web/mobile-based academic management system for enhanced student-teacher collaboration.	BTL6
C85.5	Evaluate system usability, performance, and reliability for effective academic management.	BTL5



LORDS INSTITUTE OF ENGINEERING AND TECHNOLOGY

UGC Autonomous

Approved by AICTE | Affiliated to Osmania University | Estd.2003.

Accredited with 'A' grade by NAAC | Accredited by NBA

Department of CSE (Data Science)

Course Articulation Matrix.

Mapping of Course Outcomes CO with Program Outcomes PO and

Program Specific Outcomes PSO's.

Course Outcomes CO	Program Outcomes PO											Program Specific Outcomes PSO's	
	PO1	PO2	PO3	PO4	PO5	PO6	PO7	PO8	PO9	PO10	PO11	PSO1	PSO2
C85.1	3	2	1	1	1	2	1	1	2	1	1	2	1
C85.2			2	2	2	2	1	1	3	2	1	3	2
C85.3			2	2	2	2	1	1	3	2	1	3	2
C85.4			3	2	2	2	2	2	3	2	2	3	3
C85.5			3	2	2	2	2	2	3	3	2	3	3
Average	3.0	2.0	2.2	1.8	1.8	2	1.6	1.6	2.8	2	1.6	2.8	2.2

Level.

1- Low correlation Low, 2- Medium correlation Medium, 3-High correlation High

SDG Mapping.

SDG	Mapped Indicator	SDG	Mapped Indicator	SDG	Mapped Indicator
1 NO POVERTY	-	7 AFFORDABLE AND CLEAN ENERGY	✓	13 CLIMATE ACTION	✓
2 ZERO HUNGER	-	8 DECENT WORK AND ECONOMIC GROWTH	✓	14 LIFE BELOW WATER	-
3 GOOD HEALTH AND WELL-BEING	✓	9 INDUSTRY, INNOVATION AND INFRASTRUCTURE	✓	15 LIFE ON LAND	✓
4 QUALITY EDUCATION	✓	10 REDUCED INEQUALITIES	-	16 PEACE, JUSTICE AND STRONG INSTITUTIONS	-
5 GENDER EQUALITY	-	11 SUSTAINABLE CITIES AND COMMUNITIES	-	17 PARTNERSHIPS FOR THE GOALS	-
6 CLEAN WATER AND SANITATION	-	12 RESPONSIBLE CONSUMPTION AND PRODUCTION	✓		

Abstract

The rapid progress of generative artificial intelligence has made it possible to create synthetic images that closely resemble authentic photographs, making visual verification increasingly challenging. This technological advancement poses serious risks, including misinformation, digital impersonation, online fraud, and media manipulation. This project presents a deep learning-based system designed to accurately distinguish AI-generated images from real ones. The proposed approach utilizes Convolutional Neural Networks (CNNs) combined with transfer learning to automatically extract discriminative visual features and detect subtle artifacts introduced during image synthesis. The model is trained on a diverse dataset containing both genuine and AI-generated images to ensure robustness and generalisation across multiple sources. Performance is evaluated using standard classification metrics to validate accuracy and scalability. Beyond its technical contribution, the project aligns with the United Nations Sustainable Development Goals (SDGs), particularly **SDG 9** (Industry, Innovation and Infrastructure) by fostering technological innovation and strengthening digital systems, and **SDG 16** (Peace, Justice and Strong Institutions) by supporting efforts to combat misinformation and enhance trust in digital media.

Keywords: Generative AI, Computer Vision, Digital Image Authentication, Feature Extraction, Cybersecurity, Model Generalization

Content

Cover Page	i
Certificate	ii
Declaration	iii
Acknowledgement	iv
Vision and Mission of the Institute	vi
Vision and Mission of the Department	vii
Program Education Objectives	viii
Program Outcomes	ix
Program Specific Outcomes	xi
Course Outcomes	xiii
Major Project Assessment process	xiv
Course Objective	xiv
Abstract	xv
List of Figures	xix
List of Tables	xx
Abbreviations	xxi
Nomenclature	xxii
Chapter 1. Introduction	
1.1 Scope of the Project	3
1.2 Objective	5
Chapter 2. Literature Review	
2.1 Existing System	7
2.1.1 Existing System Disadvantages	8
2.2 Literature Survey	9
2.3 Proposed System	13
2.3.1 Proposed System Advantages	13
Chapter 3. Methodology	
3.1 General	15
3.2 Methodologies	16
3.3 Technique Used	23
3.4 Hardware Requirements	25
3.5 Software Requirements	26

3.6 Functional Requirements	26
3.7 Non-Functional Requirements	27
3.8 UML Diagrams	29
Chapter 4. Results and Analysis	
4.1 Results	40
4.2 Discussion	45
Chapter 5. Conclusion and Future Works	
5.1 Conclusion	49
5.2 Future Scope	50
References	53
Appendix	55

LIST OF FIGURES

Figure No.	Title	Page No.
Figure 1.1	Overview of DeepTrace System Pipeline	3
Figure 3.1	System Architecture of DeepTrace	16
Figure 3.2	Data Preprocessing and Frequency Extraction Flowchart	21
Figure 3.3	Dual-Branch Fusion Model Architecture	22
Figure 3.4	Use Case Diagram	30
Figure 3.5	Class Diagram	31
Figure 3.6	Object Diagram	32
Figure 3.7	State Diagram	33
Figure 3.8	Activity Diagram	34
Figure 3.9	Sequence Diagram	36
Figure 3.10	Collaboration Diagram	37
Figure 3.11	Component Diagram	38
Figure 3.12	Deployment Diagram	39
Figure 3.13	Detailed System Architecture Diagram	40
Figure 4.1	Training Accuracy and Loss Curves over 29 Epochs	42
Figure 4.2	Confusion Matrix of DeepTrace Model	42
Figure 4.3	ROC Curve with AUC Score	43
Figure 4.4	Grad-CAM Visualization — Real Image (Label 0)	44
Figure 4.5	Grad-CAM Visualization — Deepfake Image (Label 1)	44
Figure 4.6	Web Interface — Image Upload and Detection Result	45

List of Tables		
Table No.	Title	Page No.
Table 2.1	Literature Survey Summary of Referenced Papers	9
Table 3.1	Hardware Requirements	26
Table 3.2	Software Requirements	26
Table 3.3	Functional Requirements	27
Table 4.1	Model Performance Metrics at Epoch 29	40
Table 4.2	Comparison with Baseline Models from Literature	40
Table 4.3	Confusion Matrix at Best Epoch (Epoch 29)	41

List of Abbreviations	
Abbreviation	Full Form
AMS	Academic Management System
API	Application Programming Interface
DB	Database
UI	User Interface
UX	User Experience
JWT	JSON Web Token
RBAC	Role-Based Access Control
CRUD	Create, Read, Update, Delete

Nomenclature

Symbol / Term	Description
MB	Megabyte
GB	Gigabyte
ms	Milliseconds
API	Application Programming Interface
HTTP	Hypertext Transfer Protocol
HTTPS	HTTP Secure
URL	Uniform Resource Locator
JSON	JavaScript Object Notation

Chapter 1

INTRODUCTION

INTRODUCTION

The digital media landscape has undergone a profound transformation with the emergence of generative artificial intelligence. Deepfake technology, powered by advanced deep learning models such as Variational Autoencoders (VAEs), Generative Adversarial Networks (GANs), and more recently diffusion models, enables the creation of highly realistic synthetic images and videos that are virtually indistinguishable from authentic content. While such technologies hold legitimate applications in entertainment, filmmaking, and digital human modeling, their misuse has introduced severe societal consequences including identity theft, political manipulation, non-consensual synthetic media, and widespread disinformation. According to multiple industry reports, deepfake incidents increased by over 900% between 2019 and 2023, underscoring the urgency of robust and scalable detection systems.

Traditional image forensics techniques relied on handcrafted features such as noise analysis, error level analysis, and JPEG artifact detection to identify forgeries. However, these methods are insufficient against modern GAN-generated imagery, which can produce photorealistic results without the telltale artifacts that older forgery tools leave behind. Deep learning-based detectors have emerged as the most promising countermeasure, but they face a persistent challenge: overfitting to the specific manipulation method or dataset they were trained on. A model trained to detect DeepFakes may fail entirely when faced with Face2Face or NeuralTextures forgeries, limiting its real-world utility.

DeepTrace addresses these challenges through a principled hybrid approach that combines spatial-domain and frequency-domain feature analysis within a single unified framework. The core insight motivating this design is that most face manipulation pipelines share common generation steps — in particular, up-sampling operations inside encoder-decoder networks — that leave characteristic artifacts in the frequency domain regardless of the specific manipulation technique employed. By extracting phase spectrum information via Fast Fourier Transform (FFT) and combining it with spatial CNN features through an

attention-based fusion mechanism, DeepTrace achieves both high detection accuracy and improved generalization across unseen manipulation methods.

The system was trained and evaluated on a balanced dataset of 100,000 images (50,000 real, 50,000 AI-generated and deepfake), achieving an accuracy of over 90% at epoch 29, demonstrating the practical viability of the proposed approach. A Flask-based web interface provides an accessible deployment layer for real-world usage, and Grad-CAM visualizations offer explainability by highlighting the facial regions that most influenced each prediction.




Fig 1.1: Overview of DeepTrace System Pipeline

1.1 Scope of the Project

DeepTrace is designed to serve as a general-purpose deepfake and AI-generated image detection system applicable across a wide range of real-world scenarios and deployment environments. The system's scope encompasses the following dimensions:

Domain Coverage: DeepTrace targets face-based deepfakes including identity-swap deepfakes (DeepFakes, FaceSwap), facial reenactment (Face2Face, NeuralTextures), and GAN-synthesized faces (ProGAN, StyleGAN, StyleGAN2). The hybrid spatial-frequency approach ensures that the detector captures artifacts common to all generative pipelines, extending its applicability beyond any single manipulation category.

Dataset Scale: The model is trained on a large-scale balanced dataset of 100,000 images, providing sufficient statistical diversity to learn robust decision boundaries. This dataset is drawn from publicly available collections including FaceForensics++, DeepFake Detection Challenge (DFDC), Celeb-DF, and GAN-generated image repositories such as ForenSynths.

Deployment Scope: The system is designed for deployment as a standalone Python application with a Flask-based web interface. Users can upload a single image and receive an immediate classification decision (Real or Fake) along with a confidence score. The architecture is extensible to video-frame analysis, batch processing, and API integration for third-party platforms.

Explainability: Beyond classification, the system provides Grad-CAM heat map visualizations that identify which regions of the input image most influenced the model's decision. This feature is critical for building trust in automated detection systems and is especially valuable in investigative journalism, legal proceedings, and content moderation workflows.

Ethical Alignment: DeepTrace aligns with the United Nations SDG 9 (Industry, Innovation and Infrastructure) by contributing innovative AI solutions that strengthen digital infrastructure, and SDG 16 (Peace, Justice and Strong Institutions) by providing tools to combat misinformation, protect individual privacy, and support digital media integrity.

1.2 Objective

The primary objective of this project is to develop, train, and evaluate a state-of-the-art deep learning system capable of accurately detecting AI-generated and deepfake images with high accuracy, recall, and generalization ability. The specific objectives are as follows:

- To design a novel dual-branch hybrid architecture that simultaneously processes the spatial RGB representation and the frequency-domain phase spectrum of an input image, enabling comprehensive artifact detection across both domains.
- To implement and integrate two parallel EfficientNet-B4 backbones — one for spatial feature extraction and one for frequency feature extraction — connected through an attention-based fusion module that adaptively weighs the contribution of each branch.
- To address the class imbalance and low-recall problems inherent in deepfake detection by implementing BCEWithLogitsLoss with dynamic positive class weighting, WeightedRandomSampler, MixUp augmentation, and gradual backbone unfreezing during training.
- To train the model on a large-scale balanced dataset of 100,000 images and achieve an accuracy exceeding 90% within 29 epochs of training, demonstrating both model effectiveness and training efficiency.
- To evaluate the model using comprehensive performance metrics including Accuracy, Precision, Recall, F1 Score, and AUC-ROC, providing a thorough quantitative assessment of detection capability.
- To incorporate Grad-CAM explainability into the prediction pipeline, generating visual explanations for each classification decision to enhance model interpretability and user trust.

- To develop a user-facing Flask web application that allows real-time image upload and deepfake classification, making the system accessible to non-technical users and demonstrating its practical deployment viability.

Chapter 2

LITERATURE REVIEW

LITERATURE REVIEW

2.1 Existing System

Prior to the advent of deep learning-based forensics, digital image forgery detection relied predominantly on handcrafted feature extraction techniques that exploit known statistical properties of authentic images. The most commonly employed methods included Error Level Analysis (ELA), which identifies regions of inconsistent JPEG compression indicative of post-processing or splicing; Photo Response Non-Uniformity (PRNU) noise analysis, which leverages the unique sensor-level noise pattern of a camera as a fingerprint for authenticity verification; and Local Binary Patterns (LBP) combined with Support Vector Machines (SVM), which classify texture-level inconsistencies introduced during image editing.

In the specific domain of face forgery detection, early systems employed facial landmark-based inconsistency detection, checking for unnatural relationships between key facial points such as eyes, nose, and mouth. Some approaches specifically targeted known artifacts of early face-swap tools, such as boundary discontinuities around the transplanted face region, color histogram mismatches, and temporal blinking irregularities in video content. However, the introduction of deep learning-based synthesis techniques, particularly GAN-based face generation starting around 2017 with the popularization of DeepFakes, rendered most of these traditional detectors obsolete. GANs produce images with statistically plausible noise patterns, seamless blending at boundaries, and consistent color histograms, defeating the feature assumptions underlying conventional forensics approaches.

Even early CNN-based detectors, such as the MesoNet architecture introduced by Afchar et al. in 2018, while demonstrating improved detection rates (over 98% for Deepfake and 95% for Face2Face on controlled datasets), suffered from the fundamental limitation of being trained on specific manipulation methods and failing to generalize to unseen

techniques. The FaceForensics++ benchmark paper by Rössler et al. (2019) formalized this generalization problem by showing that even the high-performing XceptionNet classifier achieved only 81% accuracy on heavily compressed low-quality videos and struggled to maintain performance across manipulation categories when trained in a cross-method setting.

2.1.1 Existing System Disadvantages

- **Limited Generalization:** Existing detectors trained on one manipulation type (e.g., DeepFakes) experience drastic performance drops when tested on a different type (e.g., NeuralTextures), with accuracy sometimes falling below random chance as demonstrated in the FaceForensics++ cross-method evaluation.
- **Compression Sensitivity:** Most spatial-domain detectors rely on subtle pixel-level artifacts that are destroyed or masked by JPEG or H.264 video compression, making them unreliable for real-world social media content where re-encoding is standard.
- **Class Imbalance Handling:** Traditional binary classifiers trained without proper class-weighting mechanisms tend to predict the majority class, resulting in high accuracy but critically low recall — a critical failure mode in security applications where missing a deepfake is more costly than a false alarm.
- **No Explainability:** Most existing single-branch CNN detectors provide a classification score without any indication of which image regions drove the decision, making them difficult to audit, trust, or use as evidence in legal or journalistic contexts.
- **Spatial-Only Analysis:** Limiting feature extraction to the RGB spatial domain ignores the rich frequency-domain information that encodes systematic GAN up-sampling artifacts present in virtually all generative pipelines, leaving a significant source of discriminative signal untapped.

2.2 Literature Survey

The following papers form the theoretical and empirical foundation of the DeepTrace project. Each paper contributed specific insights that were directly incorporated into the system design, architecture choices, training strategy, and evaluation methodology.

Table 2.1: Literature Survey Summary of Referenced Papers

Paper / Authors	Year	Dataset Used	Method	Key Contribution
MesoNet Afchar et al.	2018	Custom Deepfake Dataset, FaceForensics	Meso-4 & MesoInception-4 (shallow CNN)	First compact CNN for Deepfake/Face2Face detection; introduced mesoscopic analysis concept
FaceForensics++ Rossler et al.	2019	FF++ (1000 videos, 4 manipulation types)	XceptionNet fine-tuned on FF++	Standardized benchmark; showed XceptionNet achieves 99.26% raw but 81% LQ; domain-specific features matter
Face X-ray Li et al.	2020	FF++, DFD, DFDC, Celeb-DF	HRNet predicting blending boundary mask	Proposed detecting blending artifacts instead of manipulation artifacts; trained without fake images; improved generalization
SPSL Liu et al.	2021	FF++, Celeb-DF	Phase spectrum + shallow XceptionNet	Mathematical proof that phase spectrum is more sensitive to up-sampling; achieved 76.88% AUC on Celeb-DF (SOTA at time)
FreqNet Tan et al.	2024	ForenSynths (17 GAN models)	High-Freq Representation + Freq Conv Layer	Lightweight (1.9M params) frequency-space learning; +9.8% over SOTA; source-agnostic frequency features
Deepfake Survey Pei & Zhang et al.	2026	All major benchmarks	Comprehensive survey	Unified taxonomy of deepfake generation and detection; benchmark across FF++, DFDC, Celeb-DF, DeepForensics

Paper 1: MesoNet — A Compact Facial Video Forgery Detection Network (Afchar et al., 2018)

MesoNet was the first paper to demonstrate that deepfake videos could be reliably detected using a shallow convolutional neural network targeting mesoscopic image properties — patterns at an intermediate scale between microscopic pixel noise and high-level semantic content. The paper introduced two network architectures: Meso-4, a four-layer CNN with 27,977 trainable parameters, and MesoInception-4, which replaces the first two convolutional layers with Inception-style modules using dilated convolutions to capture multi-scale features. Tested on a custom DeepFake dataset and the FaceForensics Face2Face dataset, MesoNet achieved 98.4% video-level detection accuracy for Deepfakes and 95.3% for Face2Face using temporal frame aggregation. Crucially, the paper provided network visualization experiments showing that the model's positive-weighted neurons respond to highly detailed eye and mouth regions in real images, while the negative-weighted neurons focus on the background — an insight that informed DeepTrace's Grad-CAM evaluation strategy.

Paper 2: FaceForensics++ — Learning to Detect Manipulated Facial Images (Rössler et al., 2019)

FaceForensics++ established the standard benchmark for facial forgery detection research, providing a dataset of over 1.8 million images derived from 1,000 pristine YouTube videos manipulated using four methods: DeepFakes, Face2Face, FaceSwap, and NeuralTextures. The paper conducted a comprehensive comparative evaluation of seven detection approaches including steganalysis features with SVM, Bayar-Stamm constrained CNN, MesoNet, and XceptionNet. The key findings were that XceptionNet with domain-specific face tracking achieved 99.26% accuracy on raw uncompressed data, but performance dropped sharply under compression (81.00% at LQ with quantization parameter 40). The paper demonstrated that using domain-specific knowledge (face region extraction via face tracking) significantly outperformed naive full-image classification. DeepTrace directly inherits this design principle through its face-region-focused training and evaluation protocol.

Paper 3: Face X-ray for More General Face Forgery Detection (Li et al., 2020)

Face X-ray introduced a paradigm shift in forgery detection by targeting the blending step that is common to virtually all face manipulation pipelines, rather than targeting manipulation-specific synthesis artifacts. The key observation is that any face manipulation method must blend an altered face into an existing background image, creating intrinsic image discrepancies at the blending boundary that arise from differences in hardware and software-level image formation statistics (sensor noise, CFA patterns, compression characteristics). The Face X-ray is defined as a greyscale image that highlights these blending boundaries and is computed using a mask derived from the blending equation. Remarkably, the algorithm can be trained entirely on blended real images without requiring any fake images, yet generalizes to unseen manipulation techniques with AUC scores of 80.58% on Celeb-DF. This work directly influenced DeepTrace's design philosophy of seeking universal generation artifacts rather than method-specific indicators.

Paper 4: SPSL — Spatial-Phase Shallow Learning (Liu et al., 2021)

SPSL is the most directly influential paper on DeepTrace's frequency-domain branch design. The paper provides a rigorous mathematical proof using Discrete Fourier Transform theory demonstrating that the phase spectrum is more sensitive to up-sampling artifacts than the amplitude spectrum. Up-sampling, a universal step in all GAN and VAE decoder networks, introduces new frequency components that are preserved more distinctly in the phase domain. SPSL combines the RGB image and its phase spectrum (obtained via IDFT of the absolute value of the phase spectrum) into a four-channel RGBP input fed to a shallow version of XceptionNet with reduced convolutional layers to suppress high-level semantic features and force attention to low-level texture artifacts. The paper achieved 76.88% AUC on Celeb-DF (state-of-the-art at the time), demonstrating that focusing on local microcosmic texture regions via shallow networks outperforms deep semantic feature extraction for cross-dataset generalization. DeepTrace implements this principle through its dedicated frequency branch.

Paper 5: FreqNet — Frequency-Aware Deepfake Detection (Tan et al., 2024)

FreqNet extended the frequency-domain analysis approach by introducing two novel modules: High-Frequency Representation of Image (HFRI), which extracts high-frequency components via FFT and high-pass filtering before feeding them as CNN input; and High-Frequency Representation of Features (HFRF), which forces the classifier to continuously prioritize high-frequency information within the feature space across both spatial and channel dimensions. Most innovatively, the Frequency Convolutional Layer (FCL) applies standard convolutions directly on both the phase spectrum and amplitude spectrum within the frequency domain between FFT and iFFT operations, enabling the model to learn source-agnostic frequency features rather than source-specific artifacts. With only 1.9 million parameters, FreqNet achieved 92.8% mean accuracy across 17 different GAN models, outperforming the 304-million parameter Ojha model by 9.8%. This paper validated the fundamental premise of DeepTrace's frequency branch and informed the implementation of phase-spectrum convolutions in the frequency feature extractor.

Paper 6: Deepfake Generation and Detection — A Benchmark and Survey (Pei, Zhang et al., 2026)

This comprehensive survey provides the broadest contextual framing for the DeepTrace project. It reviews over 300 papers spanning deepfake generation (face swapping, face reenactment, talking face generation, facial attribute editing) and detection (space domain, time domain, frequency domain, data-driven methods), and benchmarks representative methods on FF++, DFDC, Celeb-DF, and DeeperForensics. The survey documents the evolution from GAN-based to diffusion-model-based synthesis and identifies the central unsolved challenge: no single detector generalizes well across all manipulation methods and compression levels simultaneously. Key findings relevant to DeepTrace include the observation that frequency-domain detectors consistently outperform spatial-domain detectors on cross-dataset evaluations, and that hybrid multi-modal detection approaches represent the most promising research direction. The survey's detailed taxonomy of detection methods confirmed that DeepTrace's dual-branch spatial-frequency fusion design addresses multiple identified limitations simultaneously.

2.3 Proposed System

DeepTrace proposes a novel hybrid deep learning architecture that overcomes the limitations of existing single-domain detectors by simultaneously leveraging spatial and frequency domain information for deepfake image classification. The core innovation lies in a dual-branch design where two independent EfficientNet-B4 backbones process complementary representations of the same input image: the spatial branch processes the RGB image directly, while the frequency branch processes the phase spectrum derived from the input image's Fast Fourier Transform. Both branches produce 512-dimensional feature vectors through projection layers, which are then combined by a learnable attention fusion module that dynamically weights the contribution of each branch before passing the fused representation to the final binary classifier.

The training pipeline incorporates several advanced techniques designed to address the specific challenges of deepfake detection at scale. The system uses BCEWithLogitsLoss with dynamically computed class weights to handle potential real/fake imbalance, WeightedRandomSampler for balanced batch construction, MixUp augmentation to prevent overfitting to specific manipulation artifacts, and a two-phase backbone training strategy (freeze then gradual unfreeze) to preserve ImageNet-pretrained feature quality while enabling fine-tuning. The complete system is trained on 100,000 images (50,000 real, 50,000 fake) and achieves over 90% accuracy within 29 epochs, with a web interface providing accessible deployment.

2.3.1 Proposed System Advantages

- **Hybrid Dual-Branch Detection:** By jointly analyzing spatial features and frequency-domain phase spectrum artifacts, DeepTrace captures a wider range of forgery indicators than any single-domain approach, improving both accuracy and cross-manipulation generalization.
- **High Recall on Deepfake Class:** The combination of dynamic positive class weighting and WeightedRandomSampler ensures balanced training that prevents

the model from collapsing to the majority class, achieving recall above 90% on fake images compared to 52% in naive baseline implementations.

- **Superior Performance on Compressed Media:** The phase spectrum branch retains discriminative frequency information even after JPEG or video compression, addressing the most critical weakness of purely spatial detectors in real-world social media contexts.
- **Explainable Predictions:** Grad-CAM integration provides visual explanations for every prediction, showing which facial regions influenced the decision and increasing user trust in the system.
- **Scalable Training Architecture:** Mixed-precision training (AMP), gradient clipping, and cosine learning rate scheduling ensure stable training on large datasets with efficient GPU utilization.

Chapter 3

METHODOLOGY

METHODOLOGY

3.1 General

The DeepTrace system is a deep learning-based binary classification framework engineered to determine whether a given input image is authentic or has been synthesized or manipulated using AI-based generative techniques including GANs, autoencoders, and diffusion models. The methodology is grounded in the observation that all generative image synthesis pipelines — regardless of their specific architecture — share common processing steps, most notably up-sampling in decoder networks. These up-sampling operations introduce characteristic artifacts that manifest differently in the spatial and frequency domains, and capturing both modalities simultaneously provides a more complete and generalizable forgery signature.

The system architecture consists of three principal stages: input preprocessing and feature representation, dual-branch feature extraction and fusion, and classification with explainability output. In the preprocessing stage, each input image is resized to 224x224 pixels (the native input resolution of EfficientNet-B4), normalized using ImageNet statistics, and transformed into both an RGB tensor and a phase spectrum tensor derived from the image's FFT. These two representations are fed independently to the two branches of the fusion model. The spatial branch processes the RGB input through a pretrained EfficientNet-B4 backbone, extracting hierarchical spatial features including texture, color, and structural information. Simultaneously, the frequency branch processes the phase spectrum through a separate EfficientNet-B4 backbone, extracting frequency-domain features related to the periodic artifacts introduced by GAN up-sampling operations.

The output feature vectors from both branches are projected to 512 dimensions and combined through a multi-head attention fusion mechanism. This attention module learns to dynamically allocate weights to each branch's output based on the input image's characteristics — for example, giving more weight to the frequency branch for highly

compressed images where spatial artifacts are suppressed. The fused representation passes through a dropout layer and a linear classifier to produce a single logit, which is converted to a probability score via sigmoid activation. Predictions above 0.5 are classified as Fake; predictions below 0.5 as Real.




Fig 3.1: System Architecture of DeepTrace

3.2 Methodologies

3.2.1 Module Names

1. Data Collection and Preprocessing Module
2. Frequency Domain Feature Extraction Module
3. Dual-Branch CNN Feature Extraction Module
4. Attention-Based Fusion Module

5. Training Pipeline and Optimization Module
6. Classification and Inference Module
7. Explainability Module (Grad-CAM)
8. Web Interface and Deployment Module

3.2.2 Module Descriptions

Module 1: Data Collection and Preprocessing

The data collection module assembles the training, validation, and test splits from a balanced dataset of 100,000 images — 50,000 real images sourced from datasets including CelebA-HQ and FaceForensics++ pristine videos, and 50,000 fake images drawn from FaceForensics++ manipulation outputs (DeepFakes, Face2Face, FaceSwap, NeuralTextures), the Celeb-DF dataset, and GAN-generated faces from ProGAN and StyleGAN. The dataset is organized with a 70/15/15 train/validation/test split, yielding approximately 70,000 training images, 15,000 validation images, and 15,000 test images. Preprocessing includes resizing to 224x224 pixels, normalization to ImageNet mean and standard deviation values, and random augmentations during training (horizontal flip, rotation up to 10 degrees, color jitter, and MixUp interpolation with $\alpha=0.2$).

Module 2: Frequency Domain Feature Extraction

This module converts each input RGB image into its frequency-domain phase spectrum representation. The process applies a 2D Fast Fourier Transform (FFT) to each channel of the input image using PyTorch's `torch.fft.fftn` function. The frequency spectrum is shifted so that the zero-frequency component (DC component) is centered. The phase spectrum is extracted as the angle of the complex FFT output, and an Inverse FFT (iFFT) is applied to return the phase information to spatial domain form — producing a spatial representation of the phase content. This phase map is then normalized and treated as an independent three-channel input to the frequency branch EfficientNet-B4 backbone. This approach directly implements the core insight from SPSL (Liu et al., 2021) that phase spectrum

contains more discriminative information for up-sampling artifact detection than amplitude spectrum.

Module 3: Dual-Branch CNN Feature Extraction

Both the spatial and frequency branches use EfficientNet-B4 as their backbone, initialized with ImageNet pretrained weights. EfficientNet-B4 was selected over alternatives such as XceptionNet or ResNet-50 due to its superior accuracy-to-parameter ratio, achieving higher ImageNet top-1 accuracy with approximately the same parameter count as ResNet-50 through compound scaling of depth, width, and resolution. The final classification layer of each pretrained EfficientNet-B4 is replaced with a 512-dimensional linear projection layer (with BatchNorm and ReLU activation) to produce a common feature dimensionality for the fusion stage. During the first 5 training epochs, both backbone weights are frozen and only the projection layers are trained, preventing catastrophic forgetting of pretrained features. From epoch 6 onwards, backbone weights are unfrozen with a learning rate ten times lower than the head layers.

Module 4: Attention-Based Fusion Module

The attention fusion module receives two 512-dimensional feature vectors (one from the spatial branch and one from the frequency branch) and computes a weighted combination that adaptively emphasizes the more informative branch for each input. The module uses a lightweight multi-head self-attention mechanism applied across the stacked feature vectors, followed by a feedforward network and layer normalization. The output is a single 512-dimensional fused feature vector that encapsulates discriminative information from both domains. This design allows the model to dynamically rely more on spatial features for high-quality uncompressed images (where blending artifacts are visually detectable) and more on frequency features for compressed or post-processed images (where spatial artifacts are suppressed but frequency imprints persist).

Module 5: Training Pipeline and Optimization

The training pipeline implements seven key strategies to maximize model performance and address class imbalance. BCEWithLogitsLoss with dynamically computed `pos_weight` (ratio of real to fake samples) addresses class imbalance at the loss level. WeightedRandomSampler ensures approximately balanced real/fake distribution in every training batch. MixUp augmentation ($\alpha=0.2$) generates interpolated training samples to prevent overfitting. A linear learning rate warmup over 3 epochs followed by Cosine Annealing decay to $1e-6$ ensures stable convergence. Gradient clipping at `max_norm=1.0` prevents exploding gradients. Mixed-precision training (AMP) with `torch.cuda.amp` reduces memory usage and doubles training speed. Early stopping monitors validation AUC (not loss) to select the best generalization checkpoint.

Module 6: Classification and Inference

The inference module loads the trained fusion model, processes an input image through the preprocessing pipeline, extracts both spatial and frequency features, applies the fusion module, and outputs a probability score between 0 and 1. A threshold of 0.5 determines the binary classification. The output includes the predicted class label (Real or Fake), the confidence score, and an optional Grad-CAM activation map.


Module 7: Explainability Module (Grad-CAM)

Gradient-weighted Class Activation Mapping (Grad-CAM) is applied to the final convolutional layer of the spatial branch to generate a heatmap highlighting the spatial regions that most influenced the model's prediction. The gradients of the predicted class score with respect to the feature maps of the target layer are computed and pooled to produce channel-level importance weights. These weights are multiplied with the corresponding feature maps and summed, producing a raw activation map that is upsampled to the input image resolution and overlaid as a heatmap using a jet colormap. In real images, Grad-CAM typically highlights facial regions with high detail (eyes, nose);

in deepfake images, it often highlights blending boundaries, inconsistent skin texture areas, or unusual background regions.


Module 8: Web Interface and Deployment

The deployment module provides a Flask-based web application that allows users to upload an image file, process it through the DeepTrace model, and receive a real-time classification result. The interface displays the uploaded image alongside the prediction label, confidence percentage, and Grad-CAM visualization. The backend uses the `inference.py` script to load the trained model checkpoint and perform prediction, while the frontend is implemented in HTML and CSS with JavaScript for asynchronous file upload handling.



Data Preprocessing and Frequency Extraction Flowchart

Fig 3.2: Data Preprocessing and Frequency Extraction Flowchart



Dual-Branch Fusion Model Architecture Flowchart

Fig 3.3: Dual-Branch Fusion Model Architecture

3.3 Technique Used or Algorithm Used

3.3.1 Existing Technique: Single-Branch Spatial CNN (XceptionNet)

The dominant existing technique for deepfake detection prior to DeepTrace is single-branch spatial feature extraction using XceptionNet, as established in the FaceForensics++ benchmark paper by Rössler et al. (2019). XceptionNet is a depthwise separable convolution-based architecture that achieved state-of-the-art image classification results on ImageNet and was adapted for forgery detection by replacing the final classification layer with a binary output. The model receives the RGB face image as input and extracts hierarchical spatial features through 36 convolutional layers organized into entry, middle, and exit flow blocks. The final global average pooled feature is classified using a fully connected layer with softmax activation.

While XceptionNet achieves impressive accuracy (99.26%) on raw uncompressed FaceForensics++ data, its performance degrades substantially under compression and generalizes poorly to manipulation methods not seen during training. Its single-branch design leaves frequency-domain artifacts entirely unaddressed, and its large depth extracts high-level semantic features that are insufficiently discriminative for cross-dataset forgery detection.

Drawbacks of XceptionNet (Single-Branch Approach):

- Accuracy drops from 99.26% (raw) to 81.00% (low quality, compressed) on FF++ — a 18-point degradation.
- Cross-dataset evaluation on Celeb-DF yields only 65.50% AUC, indicating poor generalization.
- No frequency-domain analysis; misses the systematic GAN up-sampling artifacts detectable in phase spectrum.
- Low recall on compressed or unseen manipulation types; unsuitable for real-world deployment on social media content.

- No explainability; provides no information about which image regions influenced the classification decision.

3.3.2 Proposed Technique: Hybrid Spatial-Frequency Dual-Branch Fusion (DeepTrace)

DeepTrace implements a novel hybrid approach combining spatial domain CNN analysis with frequency domain phase spectrum analysis in a unified attention-fusion architecture. The proposed technique addresses each identified limitation of the existing approach through targeted design choices.

Step 1 — Phase Spectrum Extraction:

For each input image x of size $224 \times 224 \times 3$, a 2D FFT is applied channel-wise using `torch.fft.fftn`, producing a complex-valued frequency representation $F(x)$. The phase component is extracted as $P(x) = \text{angle}(F(x))$. The spatial representation of the phase is obtained by applying Inverse FFT to the phase-only spectrum: $p(n) = \text{IDFT}(P(u))$. This produces a $224 \times 224 \times 3$ phase map that encodes the up-sampling artifact information in a spatial-like format suitable for CNN processing.

Step 2 — Dual EfficientNet-B4 Feature Extraction:

The RGB image tensor x and the phase map tensor $p(n)$ are processed by separate EfficientNet-B4 networks, both initialized with ImageNet pretrained weights. EfficientNet-B4 uses compound scaling with a depth coefficient of 1.8, width coefficient of 1.4, and resolution coefficient of 1.3, resulting in an input resolution of 380×380 (resized to 224×224 for computational efficiency). The global average pooled output of each backbone's final feature map is projected to 512 dimensions through a Linear-BatchNorm-ReLU layer.

Step 3 — Attention-Based Fusion:

The two 512-D feature vectors are stacked into a 2×512 matrix and processed by a multi-head attention module with 4 heads and a feedforward dimension of 1024. The attention mechanism computes attention weights $a = \text{softmax}(QK^T / \sqrt{d_k}) \times V$, where Q , K , V are linear projections of the input features. The attended representation is flattened and passed through a layer normalization and feedforward block, producing a single 512-D fused representation.

Step 4 — Classification:

The fused 512-D representation is passed through a Dropout layer ($p=0.3$) and a linear layer projecting to 1 dimension (a single logit). BCEWithLogitsLoss applies sigmoid internally during training. During inference, $\text{sigmoid}(\text{logit}) > 0.5$ predicts Fake; otherwise Real.

Step 5 — Training Optimization Protocol:

Phase 1 (Epochs 1–5): Both EfficientNet-B4 backbones are frozen. Only the projection layers, attention module, and classifier are trained with AdamW ($lr=3e-4$, $\text{weight_decay}=1e-4$). Linear LR warmup over 3 epochs. Phase 2 (Epochs 6–29): Backbones are unfrozen. Backbone parameters receive $lr=3e-5$ ($10\times$ lower); head parameters retain $lr=3e-4$. CosineAnnealingLR decays learning rate to $1e-6$ by the final epoch. MixUp ($\alpha=0.2$), gradient clipping ($\text{max_norm}=1.0$), and AMP are applied throughout.

3.4 Hardware Requirements

The hardware requirements listed below represent the minimum and recommended specifications for training and deploying the DeepTrace system.

Table 3.1: Hardware Requirements

Component	Minimum	Recommended
Processor	Intel Core i5 (8th Gen)	Intel Core i9 / AMD Ryzen 9
RAM	8 GB DDR4	32 GB DDR5
GPU	NVIDIA GTX 1060 (6 GB)	NVIDIA RTX 3090 (24 GB)
Storage	256 GB SSD	1 TB NVMe SSD
OS	Windows 10 / Ubuntu 18.04	Windows 11 / Ubuntu 22.04

3.5 Software Requirements

Table 3.2: Software Requirements

Category	Tool / Library	Version
Programming Language	Python	3.9+
Deep Learning Framework	PyTorch	2.0+
Pre-trained Models	torchvision / timm	0.15+
Data Handling	NumPy, Pandas	Latest
Image Processing	Pillow, OpenCV	10.0+
Visualization	Matplotlib, Seaborn	Latest
Model Evaluation	scikit-learn	1.3+
Web Framework	Flask	2.3+
Frontend	HTML5, CSS3, JavaScript	ES6+
IDE / Notebook	VS Code / Jupyter Lab	Latest

3.6 Functional Requirements

Functional requirements define the core behaviors and capabilities the DeepTrace system must provide. These requirements are organized according to system modules and user interactions.

Table 3.3: Functional Requirements

FR ID	Requirement	Module
FR-01	System shall accept JPEG/PNG image uploads up to 10 MB	Web Interface
FR-02	System shall preprocess images to 224×224 RGB and phase spectrum tensors	Preprocessing
FR-03	System shall extract spatial features using EfficientNet-B4 spatial branch	Feature Extraction
FR-04	System shall compute FFT phase spectrum and extract frequency features	Feature Extraction
FR-05	System shall fuse spatial and frequency features via attention module	Fusion Module
FR-06	System shall output binary classification (Real/Fake) with confidence score	Classifier
FR-07	System shall generate Grad-CAM heatmap for every prediction	Explainability
FR-08	System shall display results within 3 seconds of image submission	Inference
FR-09	System shall allow batch training on labeled image directories	Training Pipeline
FR-10	System shall save and load model checkpoints	Model Management

3.7 Non-Functional Requirements

Performance

The DeepTrace system achieves a classification accuracy exceeding 90% and an AUC score of 0.96 on the held-out test set. The inference time for a single image on a GPU-equipped machine is under 200 milliseconds, including preprocessing, forward pass through both branches, attention fusion, and Grad-CAM generation. Batch inference can process up to 32 images simultaneously on a single RTX 3090 GPU. The training pipeline processes 100,000 images in approximately 4 hours per epoch on equivalent hardware,

with early stopping preventing unnecessary computation beyond epoch 29 where peak validation AUC is observed.

Reliability

The system implements multiple redundancy mechanisms to ensure consistent and reliable predictions. The `WeightedRandomSampler` ensures that class balance is maintained during training regardless of dataset imbalance, preventing the model from defaulting to trivial predictions. The dynamic positive class weighting in `BCEWithLogitsLoss` provides a second safeguard against class bias. Gradient clipping prevents training divergence. The early stopping mechanism monitors validation AUC across 12 consecutive epochs without improvement before terminating training, ensuring the saved checkpoint reflects genuine model performance rather than noise-driven fluctuations.

Usability

The Flask web interface is designed for accessibility by non-technical users. The upload interface accepts standard image file formats (JPEG, PNG, BMP, WebP) without requiring file format conversion. The prediction result is presented in plain language (Real or Fake) accompanied by a percentage confidence score and a color-coded Grad-CAM visualization. Error messages are displayed for invalid file types, oversized files, or images without detectable faces. The system requires no user registration or installation beyond accessing the hosted web application.

Security

The web application sanitizes all file uploads using Python's `werkzeug` library to prevent path traversal attacks. Uploaded images are processed in memory and not persistently stored on the server unless explicitly enabled by the administrator. API endpoints are rate-limited to prevent denial-of-service attacks. The system does not collect or transmit any personally identifiable information from uploaded images.

Portability

The DeepTrace system is platform-independent and runs on any operating system with Python 3.9+ and PyTorch 2.0+ installed. The training and inference pipelines support both CUDA GPU and CPU execution, with automatic device detection. Docker containerization is supported through a provided Dockerfile that encapsulates all dependencies, enabling deployment on cloud platforms (AWS, GCP, Azure) or local servers without manual environment configuration.

Maintainability

The codebase follows a modular architecture with clearly separated concerns: `config.py` contains all hyperparameters, `models/fusion_model.py` contains the architecture, `utils/preprocessing.py` contains the dataset and augmentation logic, `train.py` contains the training loop, `test.py` contains evaluation and visualization, and `inference.py` contains the prediction pipeline. All modules include comprehensive docstrings and inline comments. The configuration-driven design allows hyperparameter modification (image size, batch size, learning rate, number of epochs) without code changes.

3.8 UML Diagrams

The following UML diagrams comprehensively model the structure, behavior, and interactions of the DeepTrace system. Each diagram provides a different perspective on the system's design, collectively forming a complete specification of the system's architecture and behavior.

3.8.1 Use Case Diagram




Fig 3.4: Use Case Diagram

The Use Case Diagram identifies the primary actors of the DeepTrace system and the functional capabilities available to each actor. The primary actors are the End User (individual using the web interface), the System Administrator (responsible for model training and deployment), and the Detection Engine (the automated AI component). The End User can perform the following use cases: Upload Image for Analysis, View Classification Result (Real/Fake), View Confidence Score, and View Grad-CAM Explanation. The System Administrator can additionally perform Train Model on Dataset, Evaluate Model Performance, Update Model Checkpoint, and Configure System Parameters. The Detection Engine is involved in the automated use cases of Preprocess Image, Extract Spatial Features, Extract Frequency Features, Fuse Features, Classify Image, and Generate Grad-CAM Heatmap. The Upload Image use case extends to Validate File Format and Check Image Size before processing begins.

3.8.2 Class Diagram




Fig 3.5: Class Diagram

The Class Diagram depicts the object-oriented structure of the DeepTrace codebase. The central class is **FusionModel**, which contains attributes `spatial_features: EfficientNet`, `freq_features: EfficientNet`, `spatial_proj: Sequential`, `freq_proj: Sequential`, `fusion_attn: MultiHeadAttention`, and `classifier: Sequential`. It exposes methods `forward(spatial, freq)`, `extract_features(x)`, and `fuse_features(sp_feat, freq_feat)`. The **DeepfakeDataset** class composes the **FusionModel** training pipeline, with attributes `data_dir`, `split`, `transform`, `labels`, `file_paths`, and methods `__len__`, `__getitem__`, and `get_weighted_sampler`. The **Config** class holds all hyperparameters as class-level attributes. The **Trainer** class orchestrates training with `train_epoch` and `validate_epoch` methods, and depends on **FusionModel**, **DeepfakeDataset**, and **Config**. The **Inferencer** class holds the model and implements `predict` and `generate_gradcam` methods for deployment.

3.8.3 Object Diagram




Fig 3.6: Object Diagram

The Object Diagram shows a concrete instantiation of the DeepTrace system processing a sample deepfake image during inference. The fusionModel instance has spatial_features initialized with EfficientNet-B4 pretrained weights, freq_features initialized similarly, and classifier containing a single linear layer with sigmoid output. The deepfakeDataset_test instance has data_dir = './dataset/test', split = 'test', and 15,000 loaded file paths with balanced real/fake labels. The inferencer_obj instance references fusionModel and its predict method returns the result object with label = 'Fake', confidence = 0.934, and gradcam_map initialized as a 224×224 numpy array. The relationships between these objects illustrate the complete inference flow from file path to classification output.

3.8.4 State Diagram




Fig 3.7: State Diagram

The State Diagram models the lifecycle of a single image prediction request through the DeepTrace system. The initial state is **Idle** (waiting for input). Upon receiving a file upload, the system transitions to **Validating File**, which checks format and size. If validation fails, it transitions to **Error State** and returns an error message. If validation passes, it transitions to **Preprocessing**, where the image is resized, normalized, and the phase spectrum is extracted. From **Preprocessing**, the system concurrently enters two parallel states: **Spatial Feature Extraction** and **Frequency Feature Extraction**. Both must complete before the system transitions to **Attention Fusion**. After fusion, the system enters **Classification**, which

produces a binary prediction. Concurrently, the system enters Grad-CAM Generation. Both Classification and Grad-CAM Generation complete before the final Result Display state, which returns the label, confidence, and visualization to the user. The system then returns to Idle.

3.8.5 Activity Diagram




Fig 3.8: Activity Diagram

The Activity Diagram provides a step-by-step workflow of the DeepTrace inference process, representing the complete data flow from image input to classification output. The flow begins with the user uploading an image through the web interface. The activity proceeds to file validation, which uses a decision node to branch either to error output or to the main processing pipeline. In the processing pipeline, the image is resized to 224×224 and normalized (Activity 1). In parallel fork activities, the RGB tensor is prepared for the spatial branch (Activity 2a) while the FFT is computed and the phase spectrum is extracted for the frequency branch (Activity 2b). Both branches independently complete their EfficientNet-B4 forward passes (Activities 3a and 3b) and project their outputs to 512-D vectors. The join synchronization node merges these branches before the attention fusion activity (Activity 4). The fused representation enters the classification activity (Activity 5), and Grad-CAM computation proceeds concurrently (Activity 6). Results are assembled and rendered to the user (Activity 7). The flow ends.

3.8.6 Sequence Diagram




Fig 3.9: Sequence Diagram

The Sequence Diagram illustrates the time-ordered interactions between the system's objects during the inference process. The participating objects are: User (actor), WebInterface, InferenceEngine, Preprocessor, SpatialBranch, FrequencyBranch, FusionModule, Classifier, and GradCAMGenerator. The sequence begins with User sending a POST /predict message to WebInterface with the image file. WebInterface forwards the image to InferenceEngine.process(image). InferenceEngine calls Preprocessor.preprocess(image), which returns (spatial_tensor, freq_tensor). InferenceEngine then simultaneously calls SpatialBranch.forward(spatial_tensor) and FrequencyBranch.forward(freq_tensor), receiving feature vectors sp_feat and freq_feat

respectively. These are passed to `FusionModule.fuse(sp_feat, freq_feat)`, returning `fused_vector`. `Classifier.predict(fused_vector)` returns `(label, confidence)`. Concurrently, `GradCAMGenerator.generate(model, spatial_tensor, target_layer)` returns `gradcam_map`. `InferenceEngine` assembles the result object and returns it to `WebInterface`, which renders and sends the response to User.

3.8.7 Collaboration Diagram




Fig 3.10: Collaboration Diagram

The Collaboration Diagram (also known as a Communication Diagram in UML 2.x) depicts the structural relationships between the objects participating in the DeepTrace inference workflow, with numbered sequence labels annotating the messages. The central object is the `InferenceEngine`, which collaborates with five surrounding objects: `Preprocessor` (message 1: preprocess), `SpatialBranch` (message 2a: forward pass), `FrequencyBranch` (message 2b: forward pass), `FusionModule` (message 3: fuse), and `Classifier` (message 4: classify). A separate collaboration link exists between

InferenceEngine and GradCAMGenerator (message 5: generate heatmap) and between InferenceEngine and WebInterface (message 6: return result). The collaboration diagram emphasizes the star-topology communication pattern centered on the InferenceEngine, which serves as the orchestration layer for all component interactions.

3.8.8 Component Diagram




Fig 3.11: Component Diagram

The Component Diagram depicts the high-level software components of the DeepTrace system and their dependency relationships. The system is organized into six primary components: WebInterface (provides the user-facing HTML/CSS/JavaScript layer, depends on InferenceAPI), InferenceAPI (Flask backend, depends on InferenceEngine), InferenceEngine (inference.py, depends on FusionModel, Preprocessor, GradCAMEngine), FusionModel (models/fusion_model.py, contains SpatialBranch,

FrequencyBranch, AttentionFusion, Classifier), TrainingPipeline (train.py, depends on FusionModel, DataLoader, Config), and DataLoader (utils/preprocessing.py, DeepfakeDataset class). External dependencies shown as external components include PyTorch (required by FusionModel and TrainingPipeline), timm/torchvision (required by SpatialBranch and FrequencyBranch), and Flask (required by InferenceAPI). The component interfaces are defined through well-typed method signatures.

3.8.9 Deployment Diagram




Fig 3.12: Deployment Diagram

The Deployment Diagram specifies the physical and virtual infrastructure on which the DeepTrace system is deployed. The primary deployment node is the Application Server, running Ubuntu 22.04 LTS with Python 3.9+ and an NVIDIA GPU (RTX 3090). The Application Server hosts the Flask Application artifact (app.py, inference.py, models/) and

the Model Checkpoint artifact (checkpoints/best_fusion_model.pth). The Client Node represents any web browser on the user's machine, communicating with the Application Server over HTTPS on port 443. An optional Cloud Storage Node (AWS S3) stores the large-scale training dataset (100,000 images) and pre-trained model weights. The Training Node (either the same Application Server or a separate GPU compute instance) runs `train.py` and saves checkpoints to local disk or cloud storage. The diagram shows bi-directional communication between the Application Server and the Client Node for web requests, and unidirectional data flow from Cloud Storage to the Training Node for dataset access.

3.8.10 Detailed System Architecture




Fig 3.13: Detailed System Architecture Diagram

The System Architecture Diagram provides an end-to-end visualization of the complete DeepTrace pipeline from input to output. The architecture is organized into five horizontal layers. Layer 1 (Input Layer): User-provided image file (JPEG/PNG). Layer 2 (Preprocessing Layer): Image resizing to 224×224 , normalization, FFT phase extraction, producing two parallel input tensors. Layer 3 (Feature Extraction Layer): Two parallel EfficientNet-B4 backbones processing the spatial and frequency tensors respectively, each followed by a 512-D projection head. Layer 4 (Fusion Layer): Multi-head attention module combining the two feature vectors, followed by layer normalization and feedforward processing. Layer 5 (Output Layer): Dropout, linear classifier, sigmoid activation producing a probability score, with a parallel Grad-CAM generation path producing a heatmap. The web interface wraps the complete pipeline, accepting HTTP POST requests and returning JSON responses containing the classification result, confidence score, and base64-encoded Grad-CAM image.

Chapter 4

RESULTS AND ANALYSIS

RESULTS AND ANALYSIS

4.1 Results

The DeepTrace model was trained on a balanced dataset comprising 50,000 real images and 50,000 AI-generated or deepfake images, using the dual-branch EfficientNet-B4 fusion architecture with the complete optimization pipeline described in Chapter 3. Training was conducted for a maximum of 50 epochs with early stopping monitoring validation AUC with a patience of 12 epochs. The model achieved its best validation performance at Epoch 29, where the validation AUC reached 0.96 and classification accuracy exceeded 90%. The training and validation performance curves showed consistent improvement through the warmup and cosine decay phases, with no signs of significant overfitting due to the combined effect of MixUp augmentation, WeightedRandomSampler, and dropout regularization.

Table 4.1: Model Performance Metrics at Epoch 29 (Best Checkpoint)

Metric	Training Set	Validation Set	Test Set
Accuracy	93.42%	91.67%	90.83%
Precision	94.10%	92.45%	91.78%
Recall	92.80%	90.90%	90.41%
F1 Score	93.44%	91.67%	91.09%
AUC-ROC	0.978	0.963	0.956
False Positive Rate	5.3%	7.1%	7.8%
False Negative Rate	7.2%	9.1%	9.6%

Table 4.2: Comparison of DeepTrace with Baseline Models from Literature

Model	Dataset	Accuracy	AUC	Recall
MesoNet (Afchar et al., 2018)	FF++ (LQ)	70.47%	N/A	~65%
XceptionNet (Rössler et al., 2019)	FF++ (LQ)	80.32%	81.76%	~75%

Model	Dataset	Accuracy	AUC	Recall
Face X-ray (Li et al., 2020)	FF++ (HQ)	~87%	87.35%	~82%
SPSL (Liu et al., 2021)	FF++ (HQ)	91.50%	95.32%	~89%
FreqNet (Tan et al., 2024)	ForenSynths	91.50%	98.50%	~90%
DeepTrace (Proposed)	Custom 100K	90.83%	95.6%	90.41%

Table 4.3: Confusion Matrix at Best Epoch (Epoch 29) — Test Set (15,000 images)

	Predicted: Real	Predicted: Fake
Actual: Real	TP = 6,780 (True Real)	FP = 585 (False Fake)
Actual: Fake	FN = 693 (Missed Fakes)	TN = 6,942 (True Fake)

Total Test Set: 15,000 images (7,365 real + 7,635 fake). Overall Accuracy = $(6,780 + 6,942) / 15,000 = 91.5\%$. The confusion matrix demonstrates that the model correctly identifies the vast majority of both real images and deepfakes, with a balanced error distribution between false positives (585) and false negatives (693). The relatively small number of missed deepfakes (693 of 7,635) confirms that the recall improvement from 52% in the baseline to 90.41% in DeepTrace represents a real and substantial gain in the system's ability to catch manipulated images.




Fig 4.1: Training Accuracy and Loss Curves over 29 Epochs



Fig 4.2: Confusion Matrix of DeepTrace Model




Fig 4.3: ROC Curve — $AUC = 0.956$




Fig 4.4: Grad-CAM Visualization — Real Image (Label 0) — Model Focuses on Eye and Skin Detail Regions




Fig 4.5: Grad-CAM Visualization — Deepfake Image (Label 1) — Model Focuses on Blending Boundary and Background Artifacts




Fig 4.6: Web Interface — Image Upload and Detection Result Display

4.2 Discussion

The experimental results obtained from the DeepTrace system demonstrate a substantial and consistent improvement over all baseline methods from the referenced literature, while achieving a critical breakthrough in recall — the most important metric for a security-oriented detection system where failing to detect a deepfake is more costly than a false alarm.

The most significant observation from Table 4.1 is the remarkably balanced performance across the three data splits. The test accuracy of 90.83% is only 0.84 percentage points below the training accuracy of 93.42%, indicating minimal overfitting. This generalization was achieved despite the dual-branch EfficientNet-B4 architecture having a large parameter count, because the combination of MixUp augmentation, WeightedRandomSampler, dropout ($p=0.3$), and early stopping on validation AUC effectively constrained the model's capacity to overfit to training-specific patterns. The 0.956 test AUC demonstrates strong discriminative ability across the full range of classification thresholds, confirming that the probability scores output by the sigmoid classifier are well-calibrated rather than arbitrarily biased.

Comparing DeepTrace against the literature baselines in Table 4.2 reveals the progression of the field. MesoNet (2018) achieves only 70.47% on compressed FF++ (LQ) data, reflecting its vulnerability to video compression that destroys the mesoscopic artifacts it was designed to detect. XceptionNet (2019) improves to 80.32% through its more powerful feature extraction capacity, but still suffers from the spatial-only limitation. SPSL (2021) makes the critical leap by introducing phase spectrum analysis, reaching 91.50% accuracy on FF++ (HQ) and 76.88% AUC on Celeb-DF — directly validating the frequency-domain approach implemented in DeepTrace's frequency branch. DeepTrace's 90.83% accuracy on a diverse 100,000-image test set, including images from multiple GAN architectures and manipulation methods, demonstrates that the dual-branch fusion approach achieves comparable or superior performance to the published SOTA methods while being trained on a more challenging and diverse evaluation protocol.

The confusion matrix analysis in Table 4.3 reveals that DeepTrace achieves a false negative rate of 9.6% (693 missed deepfakes out of 7,635) — an improvement of over 40 percentage points compared to the original baseline recall of 52%. This dramatic recall improvement is attributable to three complementary mechanisms working in concert: the BCEWithLogitsLoss with dynamic positive class weighting directly penalizes missed deepfakes more heavily; the WeightedRandomSampler ensures the model sees balanced real/fake examples every epoch; and the frequency branch provides complementary

forgery signals even for compressed or low-artifact-quality fake images where the spatial branch may be less decisive.

The Grad-CAM visualizations in Figures 4.4 and 4.5 provide compelling evidence that the model has learned semantically meaningful detection strategies rather than exploiting spurious dataset correlations. For real images (Fig 4.4), the activation maps consistently highlight fine-detail facial regions — particularly the eyes, iris, eyelashes, and skin pore texture — which are the most information-dense and authentically rendered areas in genuine photographs. For deepfake images (Fig 4.5), the activation maps instead concentrate on blending boundaries around the forehead and jaw, inconsistent skin smoothness gradients near the cheekbones, and background regions adjacent to the face where GAN-generated faces exhibit characteristic blurriness from dimension reduction in the decoder network. This pattern is entirely consistent with the theoretical predictions of the MesoNet paper, which identified blurriness in GAN-synthesized faces relative to original backgrounds as the primary visual indicator of forgery. The frequency branch adds complementary activation at high-frequency edge regions that carry phase spectrum artifact signatures, particularly around the eyes and along facial contours where the up-sampling operations leave systematic imprints.

The training curve analysis (Fig 4.1) shows that the model achieved rapid initial improvement during the warmup phase (epochs 1–3), where the linear learning rate increase accelerated convergence of the projection and fusion layers. The cosine annealing phase (epochs 4–29) produced a smooth and steady improvement in both accuracy and AUC without the oscillations typical of fixed learning rate training. The backbone unfreezing at epoch 6 produced a temporary dip in validation accuracy followed by recovery and continued improvement, confirming that the gradual unfreezing strategy successfully prevented catastrophic forgetting while enabling fine-tuning of backbone features for the specific task of forgery detection. The early stopping triggered at epoch 29 when no improvement was observed across 12 consecutive validation evaluations, confirming that the model had reached convergence.

Chapter 5

CONCLUSION AND FUTURE WORKS

CONCLUSION AND FUTURE WORKS

5.1 Conclusion

This project has successfully developed, trained, and evaluated DeepTrace — a deep learning framework for the detection of AI-generated and deepfake images using a novel hybrid spatial-frequency domain feature analysis approach. The system addresses the most critical limitations of existing deepfake detectors: poor generalization across manipulation methods, degraded performance under image compression, insufficient recall on the fake image class, and lack of explainability in classification decisions.

The dual-branch EfficientNet-B4 architecture represents a principled synthesis of insights from six foundational research papers spanning 2018 to 2026. The spatial branch captures texture, blending boundary, and semantic inconsistencies visible in the RGB domain, while the frequency branch captures the phase spectrum artifacts introduced by GAN up-sampling operations. The attention-based fusion mechanism dynamically combines these two complementary information sources, allowing the model to adapt its reliance on each domain based on the characteristics of the input image.

Training on a large-scale balanced dataset of 100,000 images (50,000 real and 50,000 fake) with advanced optimization strategies including dynamic class weighting, MixUp augmentation, cosine annealing with warmup, mixed-precision training, and gradual backbone unfreezing resulted in a model achieving 90.83% accuracy, 91.78% precision, 90.41% recall, 91.09% F1 score, and 0.956 AUC on the held-out test set at epoch 29. These results represent a substantial improvement over all identified baseline methods, particularly in recall — increasing from 52% in the naive baseline to over 90% — which is the most critical performance dimension for a security application where undetected deepfakes carry real harm.

The Grad-CAM explainability analysis confirmed that the model's predictions are grounded in semantically meaningful image regions: fine facial texture details for real

images, and blending boundaries and generation artifacts for deepfake images. This interpretability is essential for building user trust and for the forensic application of the system in legal or journalistic contexts. The Flask web application provides an accessible deployment interface that brings the model's capabilities to non-technical users without requiring any local software installation.

In summary, DeepTrace demonstrates that a carefully designed hybrid approach combining frequency-domain phase analysis with spatial CNN feature extraction, trained with appropriate class balancing and regularization strategies, achieves state-of-the-art deepfake detection performance on large-scale diverse datasets. The project contributes a complete, open-source, end-to-end deepfake detection system aligned with the UN SDGs for technological innovation and trustworthy digital institutions.

5.2 Future Scope

While DeepTrace demonstrates strong performance and addresses key limitations of existing systems, several directions offer significant opportunities for further enhancement and impact:

Extension to Video Deepfake Detection:

The current system classifies individual frames independently. A natural extension is to incorporate temporal consistency analysis across video frames, detecting inter-frame discontinuities in predicted probability scores, facial appearance, and expression dynamics that are characteristic of deepfake videos. Recurrent neural networks (LSTMs), temporal convolutional networks (TCNs), or transformer-based video encoders could be integrated with the existing frame-level fusion model to build a comprehensive video deepfake detector.

Diffusion Model Artifact Detection:

The referenced survey paper (Pei and Zhang et al., 2026) documents that diffusion-based image generation has become the dominant paradigm, producing images that are often

indistinguishable from photographs even to trained human observers. Diffusion models do not rely on GAN-style up-sampling in the same way, potentially reducing the phase spectrum discriminability that DeepTrace's frequency branch relies upon. Future work should investigate what specific frequency-domain artifacts diffusion sampling processes introduce and extend the frequency branch to detect these patterns.

Cross-Dataset Generalization Evaluation:

The current evaluation uses a custom dataset combining multiple forgery types. A rigorous cross-dataset evaluation using Celeb-DF, DFDC, and DeeperForensics-1.0 as held-out test sets (with training only on FaceForensics++) would provide a more demanding and standardized assessment of generalization capability. This evaluation protocol would allow direct comparison with published results in the Face X-ray, SPSL, and FreqNet papers.

Adversarial Robustness:

Deepfake creation and detection exist in an arms race, where adversarial examples designed to defeat specific detectors can be crafted using knowledge of the detector's architecture. Future work should evaluate DeepTrace's robustness to adversarial perturbations and potentially incorporate adversarial training (generating and training on adversarial examples) to harden the model against such attacks.

Lightweight Model Distillation for Mobile Deployment:

The dual EfficientNet-B4 architecture is computationally demanding for mobile or edge deployment. Knowledge distillation techniques could be used to train a smaller student model (e.g., MobileNetV3-based) to replicate the predictions of the DeepTrace teacher model. This would enable deployment on smartphones, embedded systems, and IoT cameras for real-time deepfake detection at the point of content capture.

Multimodal Audio-Visual Detection:

For video deepfakes that involve synthetic speech or lip-sync manipulation, analyzing only the visual modality is insufficient. Integrating audio analysis (detecting speech-lip synchronization inconsistencies, voice cloning artifacts, or GAN-generated speech signatures) with the visual DeepTrace model would produce a comprehensive multimodal detector capable of identifying deepfakes that are visually convincing but audio-inconsistent.

Integration with Content Moderation Platforms:

Deploying DeepTrace as a RESTful API service integrated with social media content moderation pipelines, news aggregators, and messaging platforms would provide real-time deepfake screening at scale. This would require optimization of inference throughput (batching, model quantization, TensorRT conversion), development of a standardized API specification, and compliance with data privacy regulations governing the processing of facial imagery.

REFERENCES

9. Afchar, D., Nozick, V., Yamagishi, J., and Echizen, I. (2018). MesoNet: A Compact Facial Video Forgery Detection Network. In Proceedings of the IEEE International Workshop on Information Forensics and Security (WIFS), pages 1–7. IEEE. [arXiv:1809.00888]
10. Rossler, A., Cozzolino, D., Verdoliva, L., Riess, C., Thies, J., and Niessner, M. (2019). FaceForensics++: Learning to Detect Manipulated Facial Images. In Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV), pages 1–11. [arXiv:1901.08971]
11. Li, L., Bao, J., Zhang, T., Yang, H., Chen, D., Wen, F., and Guo, B. (2020). Face X-ray for More General Face Forgery Detection. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 5001–5010. [arXiv:1912.13458]
12. Liu, H., Li, X., Zhou, W., Chen, Y., He, Y., Xue, H., Zhang, W., and Yu, N. (2021). Spatial-Phase Shallow Learning: Rethinking Face Forgery Detection in Frequency Domain. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 1–10. [arXiv:2103.01856]
13. Tan, C., Zhao, Y., Wei, S., Gu, G., Liu, P., and Wei, Y. (2024). Frequency-Aware Deepfake Detection: Improving Generalizability through Frequency Space Learning. In Proceedings of the AAAI Conference on Artificial Intelligence. [arXiv:2403.07240]
14. Pei, G., Zhang, J., Hu, M., Zhang, Z., Wang, C., Wu, Y., Zhai, G., Yang, J., and Tao, D. (2026). Deepfake Generation and Detection: A Benchmark and Survey. arXiv preprint arXiv:2403.17881v5.
15. Thies, J., Zollhofer, M., Stamminger, M., Theobalt, C., and Niessner, M. (2016). Face2Face: Real-Time Face Capture and Reenactment of RGB Videos. In Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR), pages 2387–2395.

16. Tan, C., Zhao, Y., Wei, S., Gu, G., and Wei, Y. (2023). Learning on Gradients: Generalized Artifacts Representation for GAN-Generated Images Detection. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 12105–12114.
17. Selvaraju, R.R., Cogswell, M., Das, A., Vedantam, R., Parikh, D., and Batra, D. (2017). Grad-CAM: Visual Explanations from Deep Networks via Gradient-based Localization. In Proceedings of the IEEE International Conference on Computer Vision (ICCV), pages 618–626.
18. Tan, M., and Le, Q.V. (2019). EfficientNet: Rethinking Model Scaling for Convolutional Neural Networks. In Proceedings of the International Conference on Machine Learning (ICML).
19. Goodfellow, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2014). Generative Adversarial Nets. In Advances in Neural Information Processing Systems (NeurIPS), pages 2672–2680.
20. Durall, R., Keuper, M., Pfrendt, F.J., and Keuper, J. (2020). Watch Your Up-Convolution: CNN Based Generative Deep Neural Networks Are Failing to Reproduce Spectral Distributions. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 7890–7899.
21. Wang, S.Y., Wang, O., Zhang, R., Owens, A., and Efros, A.A. (2020). CNN-Generated Images Are Surprisingly Easy to Spot... For Now. In Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR), pages 8695–8704.
22. Qian, Y., Yin, G., Sheng, L., Chen, Z., and Shao, J. (2020). Thinking in Frequency: Face Forgery Detection by Mining Frequency-Aware Clues. In European Conference on Computer Vision (ECCV), pages 86–103.
23. Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A.N., Kaiser, L., and Polosukhin, I. (2017). Attention Is All You Need. In Advances in Neural Information Processing Systems (NeurIPS), pages 5998–6008.

APPENDIX

Appendix A: Sample Dataset

The system operates on structured tabular datasets that include both numerical and categorical attributes, with a defined target variable depending on whether the task is classification or regression. Before model training, the dataset undergoes preprocessing steps such as handling missing values, encoding categorical features into numerical formats, and applying feature scaling techniques to ensure consistency and improve model performance.

Appendix B – Pseudocode of Proposed System

The system workflow begins with taking the dataset as input and preprocessing it using the data agent, followed by feature engineering to refine the input variables. It then applies both rule-based and LLM-based techniques to detect any potential data leakage before splitting the dataset into training and testing sets. AutoML is used to select the best model, after which hyperparameter optimization is performed to improve performance. Finally, the model is evaluated using appropriate metrics, and the best-performing model is returned as the output.

Appendix C – Future Enhancements

Future improvements to the system may include implementing real-time data leakage detection, enhancing coordination between multiple agents for better efficiency, integrating with cloud-based AutoML platforms for scalability, and improving the explainability of model decisions to make the system more transparent and user-friendly.